

Today's plan

We will discuss divide and conquer with two problems

1. Inversion counting

2. Closest pair (IA1)

Divide and Conquer

- General approach:
 - Break problem into smaller sub-problems
 - Solve smaller sub-problems via recursion
 - Combine solutions of sub-problems to get a solution to the original problem
- A special case of recursion
- Reduce a given problem to multiple smaller instances of the original problem
 - Typically constant factor smaller ($n \rightarrow n/b$)

Problem: Inversion counting

- Input: an array A containing numbers 1,2,...,n in arbitrary order
- Output: number of inversions, i.e., the number of pairs (i, j) such that $i < j$ and $A[i] > A[j]$

Example

- Input (1, 4, 2, 5, 3)

Why study this problem?

- Consider a set of n movies

You and I can both rank them according to how much we like them

Mine: 1, 2, 3, 4, ... n

Yours: 5, 4, 1, ...

Measure the difference between two ranked lists --- a fundamental operation behind all the recommender systems (collaborative filtering)

High level idea

- Brute force counting
- Divide and conquer
 - Break the array into two parts
 - Count the left part
 - Count the right part
 - Count in the inversions cross the two parts

High level algorithm

Count(A, n)

if $n=1$ return 0

else

$x = \text{Count}\left(A_l, \frac{n}{2}\right)$

$y = \text{Count}\left(A_r, \frac{n}{2}\right)$

$z = \text{CountCross}(A_l, A_r)$

return $x+y+z$

countcross(A_l, A_r) needs to run in $O(n)$ in order to achieve an overall runtime of $O(n \log n)$

What if A_l and A_r are already sorted?

Can we count the cross inversions in $O(n)$ time?

$$A_l = (2,4,5) \quad A_r = (1,3,6)$$

Building on Merge_sort

Sort-and-Count(A, n)

if $n=1$ return 0

else

$$(L_s, x) = \text{Sort-and-Count}\left(A_l, \frac{n}{2}\right)$$
$$(R_s, y) = \text{Sort-and-Count}\left(A_r, \frac{n}{2}\right)$$
$$(A_s, z) = \text{Merge-and-CountCross}(L_s, R_s)$$

return $(A_s, x + y + z)$

Pseudo code for merge

L : left-half sorted array (size n)

R : right-half sorted array (size m)

A : output sorted merged array

$i = 0; j = 0$

while $i < n$ and $j < m$

 if $L(i) \leq R(j)$

$A(k) = L(i); i++$

 else

$A(k) = R(j); j++$

 end if

End while

(Ignore the end case)

What would happen if there is no inversion between L and R ?

How many inversions can we be sure of when the else branch is taken?

Example

- Consider merging $(1, 3, 5)$ and $(2, 4, 6)$

Pseudo code for Merge-and-CountCross

```
L: left-half sorted array (size n)
R: right-half sorted array (size m)
A: output sorted merged array
z: the # of cross inversions


---


i = 0; j = 0; z = 0
While i < n and j < m
    if L(i) ≤ R(j)
        A(k) = L(i); i ++
    else
        A(k) = R(j); j ++
        z = z + (n - i)
    end if
End while
                                (Ignore the end case)
```

Run time of subroutine:
 $O(n)$

Run time of the overall algorithm:

$$T(n) = 2T\left(\frac{n}{2}\right) + O(n)$$

Same as merge_sort:
 $O(n \log n)$

2nd Problem: Closest Pair of Points

Problem Definition:

Given n points in 2-d plane, find a pair or pairs with the smallest Euclidean distance between them.

Our goal:

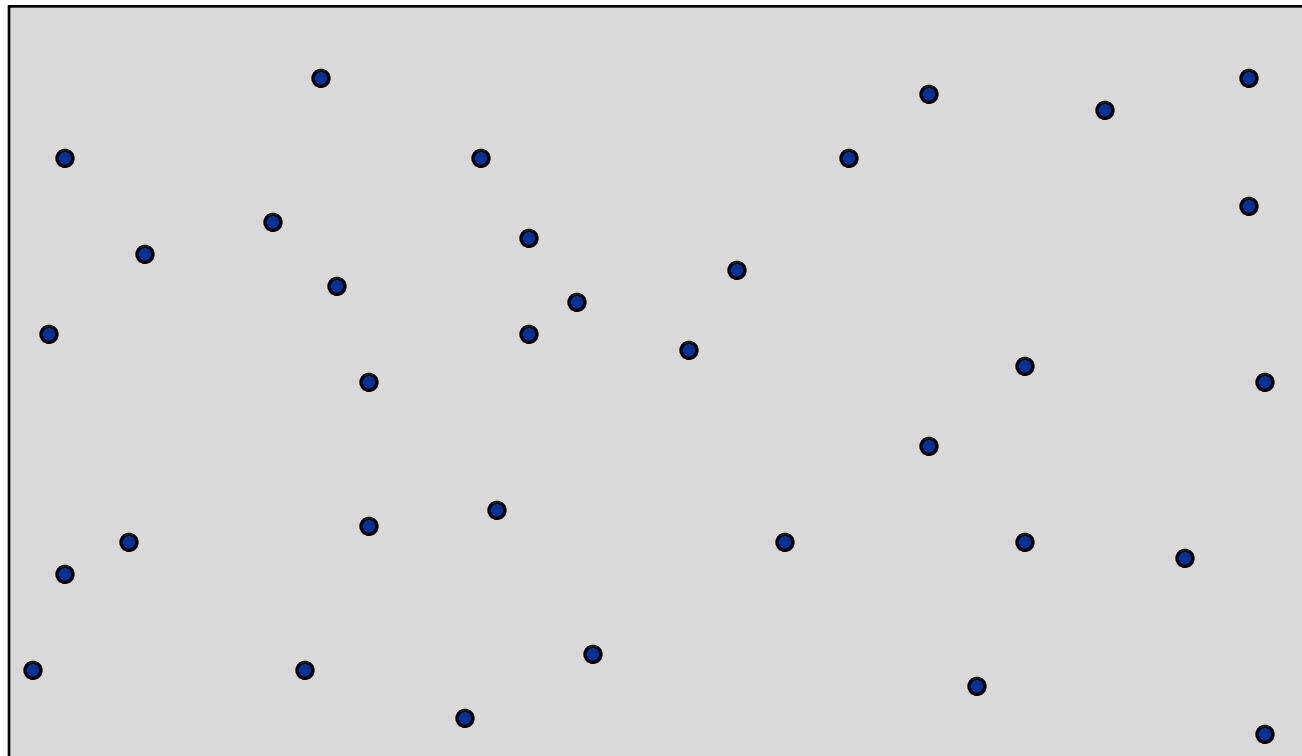
Design a divide and conquer algorithm that runs in $O(n \log n)$

Preliminary thoughts:

- Brute force – directly compute the distances between all pairs and find and record the pairs with the smallest distance
- If 1-d, we have an easy $n \log n$ algorithm by sorting the points and scan the distances between neighboring points.

Let's try to design a divide and conquer algorithm

- How to divide?



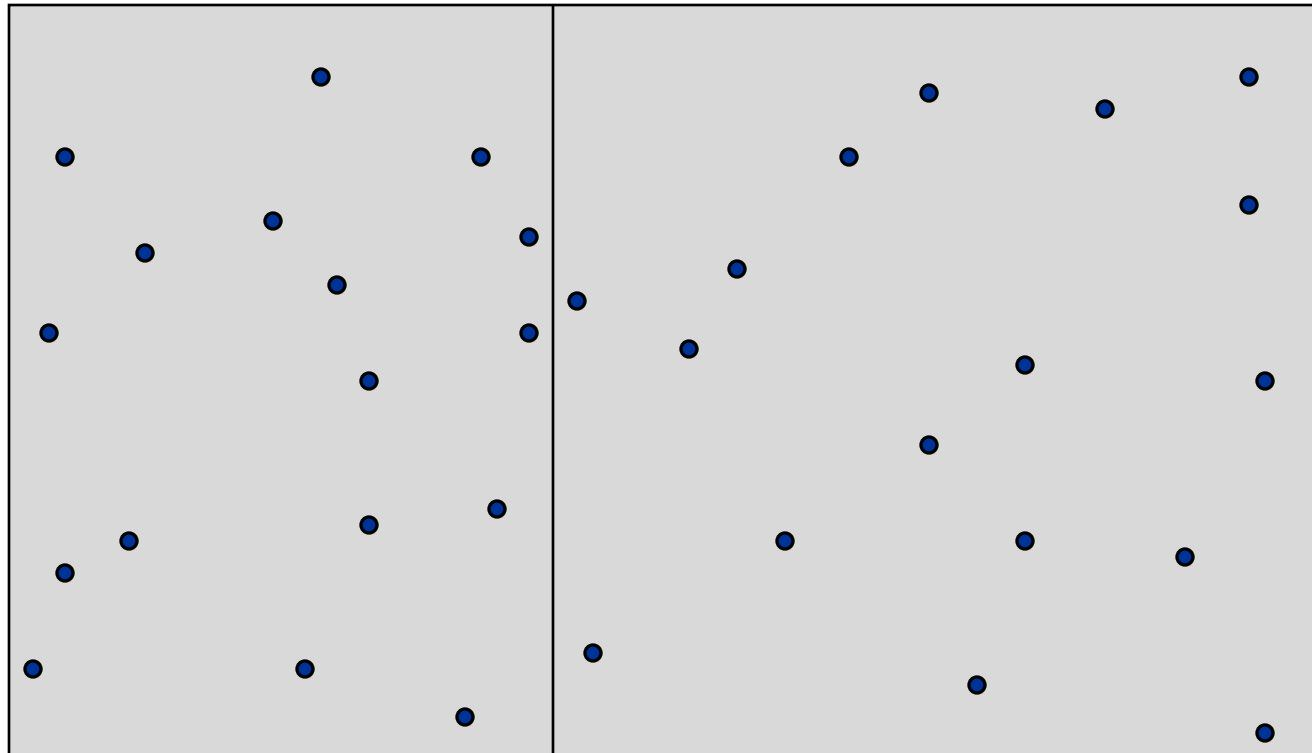
High level idea

Divide: draw vertical line s.t. $\sim n/2$ points on each side.

Conquer: find closest pair in each side recursively.

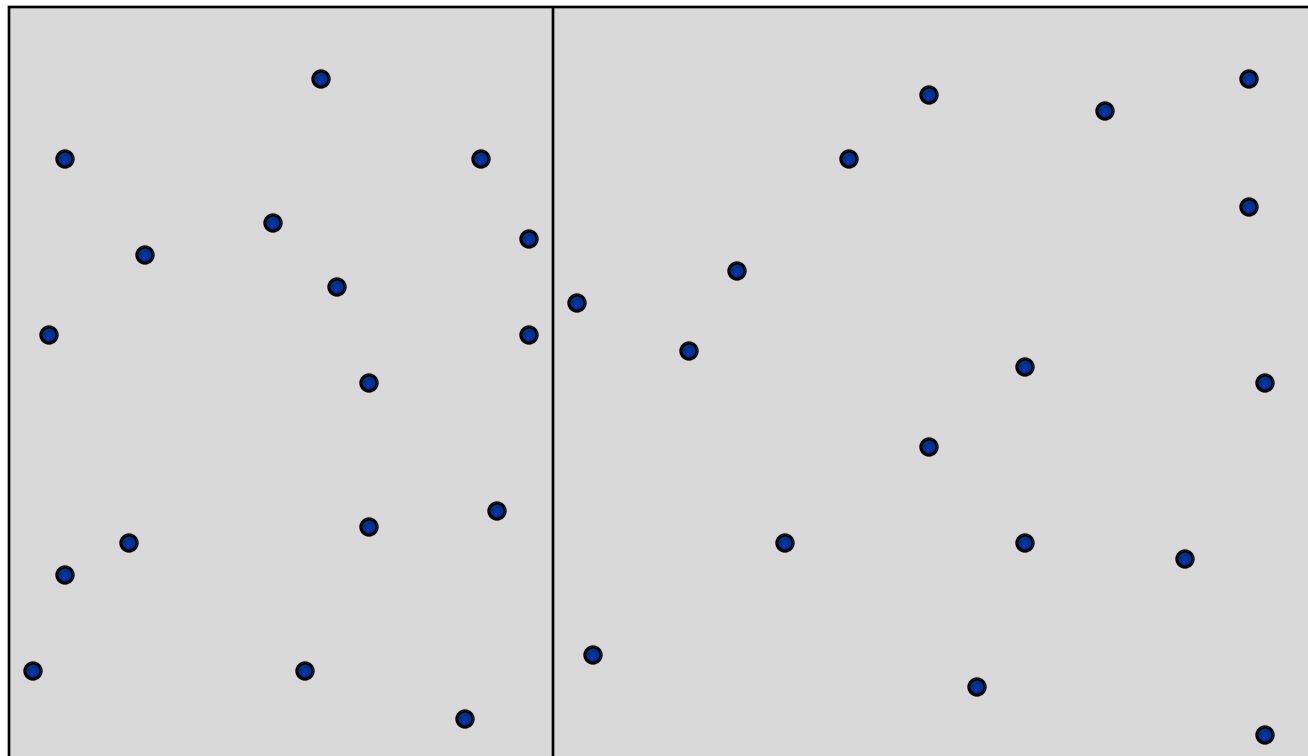
Combine: find closest pair with one point in each side.

Return best of 3 solutions.



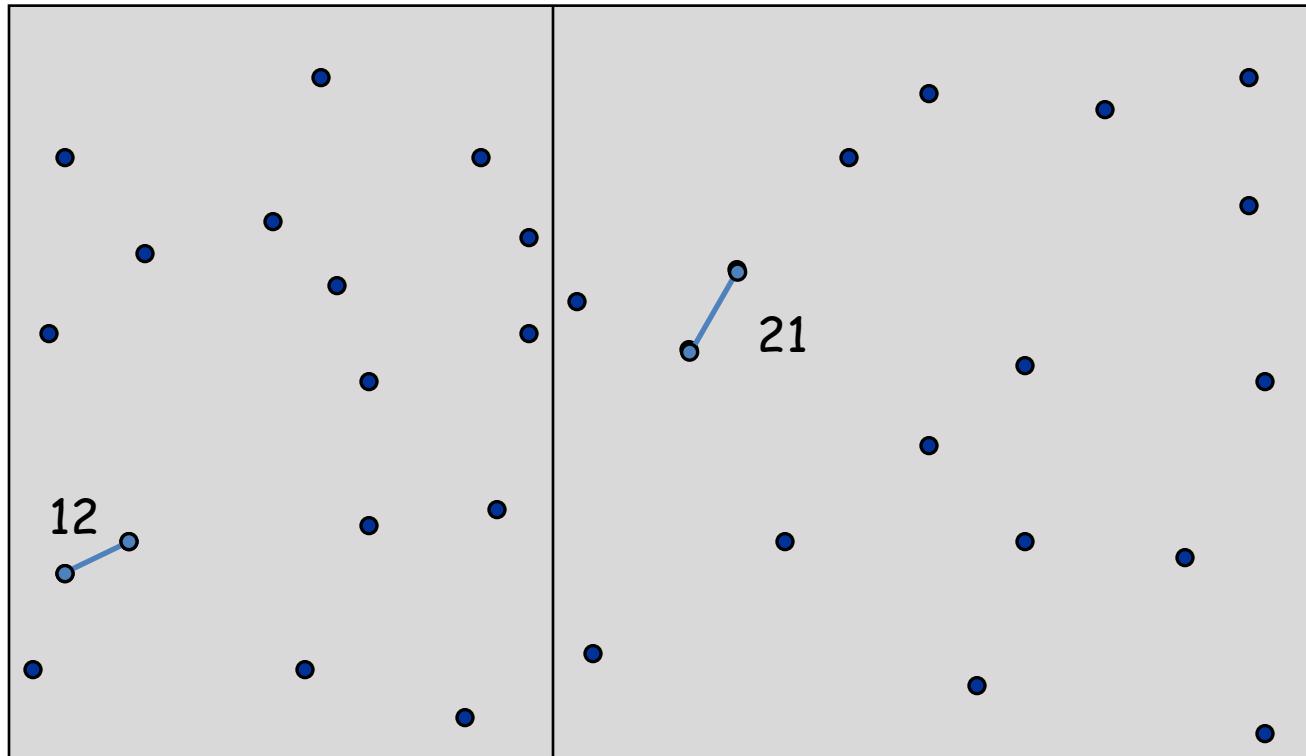
Finding the closest pair across

- Naïve approach: consider every pair across

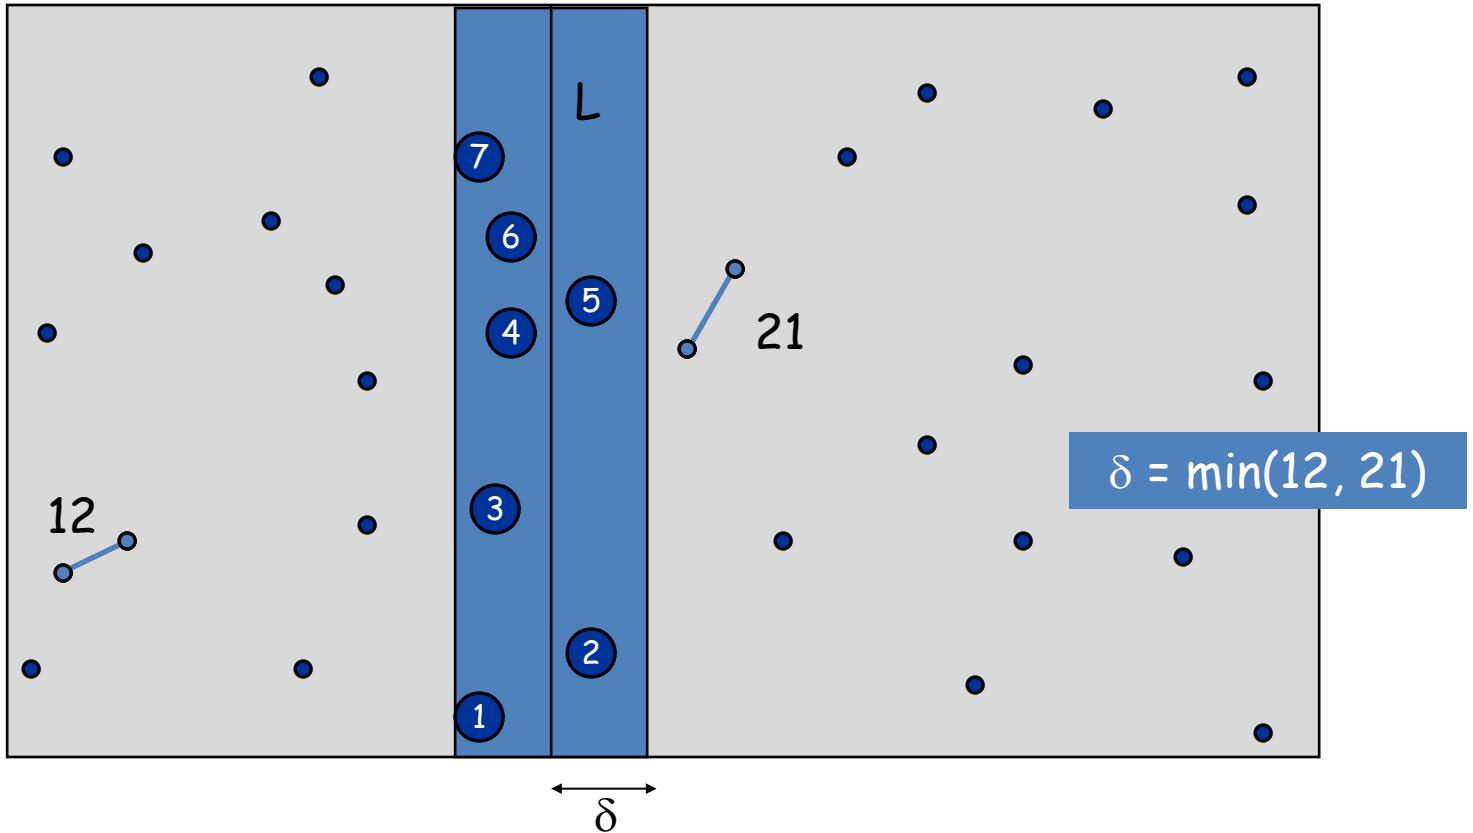


Getting smarter

- If the shortest so far is δ , then no need to consider pairs more than δ apart
- No need to consider points δ away from the median line



Focus on the middle strip



- Issue: in the worst case $O(n)$ points will be in the middle strip
- We need more pruning!!!
 - Again, we do not need consider anything more than δ apart (use the y axis to prune)

Scan the middle strip for cross pairs

Closest-cross-pairs(M_y, δ)

$M_y = \{p_1, p_2, \dots, p_m\}$: the list of points in the middle strip sorted in increasing order of y

1. $d_m = \delta$
2. for $i = 1$ to $m - 1$
3. $j = i + 1$
4. while $p_j(y) - p_i(y) \leq \delta$ and $j \leq m$
5. $d = D(p_i, p_j)$
6. $d_m = \min\{d, d_m\}$
7. $j = j + 1$
8. end while
9. end for
10. Return d_m

Claim:

For any pair (i, j) , if $D(i, j) \leq \delta$, we must have $p_j(y) - p_i(y) \leq \delta$, thus it will be compared with line 5,6!

But this has double loop? Could it be linear?

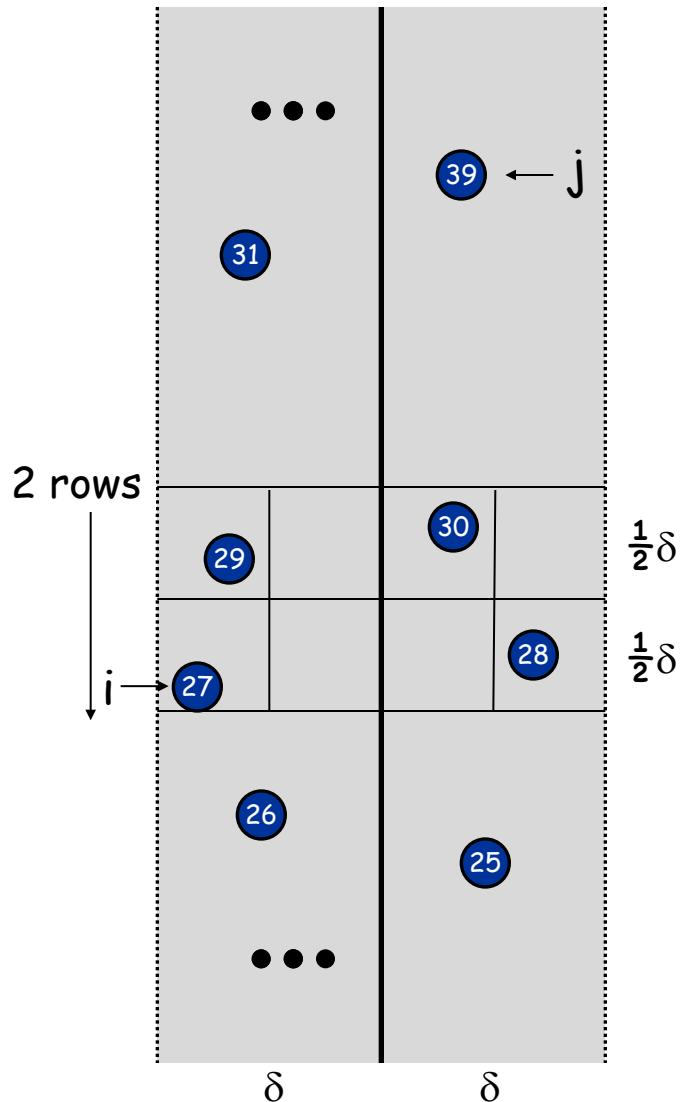
Claim. For any point p_i , the while loop will execute at most 7 times

Proof:

Consider all points above p_i that have y value within δ of $p_i(y)$

Together with p_i they must lie in side the rectangle of height δ and width 2δ

The rectangle can be divided into 8 cells. There are at most 1 point inside each cell.



Closest Pair Algorithm

(this does not keep track of all the pairs, needs to be added)

```
Closest-Pair( $p_1, \dots, p_n$ ) {  
1. if  $n \leq 3$   
2.   compute and return the min distance  
3. else  
4.   Compute separation line  $L$   
5.    $\delta_1 = \text{Closest-Pair}(\text{left half})$   
6.    $\delta_2 = \text{Closest-Pair}(\text{right half})$   
7.    $\delta = \min(\delta_1, \delta_2)$   
  
8.   Identify all points within  $\delta$  from  $L$   
9.   Sort them by y-coordinate into  $M_y$   
10.   $d_m = \text{closest-cross-pair } (M_y, \delta)$   
  
11.  return  $d_m$ .  
}
```

$O(n \log n)$

$2T(n/2)$

$O(n)$

$O(n \log n)$

$O(n)$

Enhanced Divide and Conquer

- Pre-sort all points based on x and y coordinates respectively
- For Line 8-9, scan the master list pre-sorted based on y to create M_y by excluding points δ away from L in x -coordinate